

VISHWAS PATEL

669-340-8432 ◇ Santa Clara, CA ◇ patelvishwas2401@gmail.com ◇ [LinkedIn](#) ◇ [Portfolio](#) ◇ [GitHub](#)

EDUCATION

Master of Science in Computer Science

Santa Clara University, GPA: 3.56/4

Sept. 2024 – June 2026

Santa Clara, CA

Bachelor of Engineering in Computer Engineering

Gujarat Technological University, GPA: 3.75/4

July 2019 – May 2023

Gandhinagar, India

SKILLS

Languages	Python, Java, C/C++, GO, JavaScript, TypeScript, Bash, SQL
Backend / Web	FastAPI, Flask, Django, React, Node.js, REST APIs, SQLAlchemy
Databases	Snowflake, PostgreSQL, MySQL, MongoDB
AI / Data	TensorFlow, Google ASR, OpenAI Whisper, ElevenLabs, Apache Spark, Apache Airflow
DevOps / Tools	Docker, GCP, Redis, GitHub Actions (CI/CD), Grafana, Git

EXPERIENCE

Data Engineering – Student Worker

Santa Clara University | Python, Snowflake, REST APIs, Grafana

Sept. 2025 – Present

Santa Clara, CA

- Developed a Snowflake ingestion pipeline processing \$350K+ weekly POS transactions, enabling real-time analytics for price optimization and demand forecasting.
- Centralized observability for 500+ IoT devices (door locks, credential readers, POS) via Grafana Cloud, reducing mean-time-to-detect issues by 60% and enabling early fault detection across the fleet.
- Boosted downstream query performance by 40% by implementing a Redis caching layer and optimized Snowflake indexing for high-frequency REST APIs.

Software Engineer

Shaip | FastAPI, Python, ElevenLabs, Google ASR, OpenAI Whisper, GCP

Aug. 2023 – May 2024

Ahmedabad, India

- Powered 300+ daily multilingual audio transcriptions across 5+ languages, cutting per-transcription cost by 30%, by architecting a FastAPI routing engine across Google ASR, Whisper, and ElevenLabs.
- Enabled fail-safe ingestion for audio files up to 2GB by designing a CI/CD pipeline to build and deploy auto-scaled GCP microservices with CRIU checkpointing, eliminating full-job reprocessing on failure.

PROJECTS

RedWatch: SOC Alert Detection Platform

[GitHub](#)

- Engineered a full-stack threat detection platform identifying 8 MITRE-aligned attack patterns (C2 beaconing, SQL injection, data exfiltration), enabling rapid SOC triage for high-volume security logs.
- Optimized LLM inference pipeline via adaptive log sampling and payload compaction, reducing token consumption by 55% while maintaining structured, audit-ready incident report quality.
- Automated full-stack deployment pipelines via GitHub Actions for a containerized React/TS, Flask, and PostgreSQL architecture, integrating VirusTotal API threat intelligence across 70+ AV engines.

Sentinel: Incident Response Automation System

[GitHub](#)

- Achieved 95% root cause confidence by implementing a 4-agent pipeline (Classifier → Investigator → Hypothesis → Responder) with hybrid RAG retrieval (60% FAISS + 40% BM25) using Anthropic Claude.
- Engineered a distributed, Redis-backed task queue to orchestrate concurrent multi-agent investigations, implementing state management, real-time observability, and graceful failure recovery mechanisms.
- Reduced end-to-end incident analysis latency by 35% by parallelizing independent FAISS and BM25 retrieval paths via Python's asyncio event loop, eliminating synchronous I/O blocking.